



Learner Guide

Generative AI for Social Media Marketing

C1468 | | 2 days / 15 hours | v1.0 | 13 Sep 2026

Tertiary Infotech Academy Pte Ltd • UEN 201200696W

Document Version Control Record

Version	Date	Material change	Owner
Legacy v10	Prior release	Instagram-focused source deck; retained as reference	Tertiary Infotech
v1.0	13 Sep 2026	First AI social marketing rebuild with 201 slides and eight labs	Tertiary Infotech
v1.0	13 Sep 2026	Added admin visuals and document navigation	Tertiary Infotech

Table of Contents

Course identity and outcomes • delivery topics • detailed lab workflow • references

Course identity and outcomes

Official course page: <https://www.tertiarycourses.com.sg/generative-ai-for-social-media-marketing.html>

- LO1: Develop a plan and objectives for a social marketing campaign
- LO2: Evaluate opportunities and the feasibility of social media platforms
- LO3: Create a generative AI social media campaign and monitor its outcome
- LO4: Evaluate AI campaign use and organisational impact

How to use this guide

Follow the eight lab folders in order. Each folder contains a self-contained brief, mock data, editable marketing asset, a prompt PDF and an evidence checklist. Use synthetic data only. Obtain approval before publishing anything to a live account.

Accessing the course materials

1. Open <https://lms-tms.tertiaryinfotech.com/> in a browser and sign in with the learner account supplied by the provider.
2. Find the course record titled Generative AI for Social Media Marketing and confirm course code C1468.
3. Open the courseware area. Download the Learner Slides PDF and Learner Guide PDF, checking that each is the current v1.0 copy.
4. Open the Activities/Lab folder, then open the numbered Lab-01 through Lab-08 folders in order. Each contains README.pdf, prompts.pdf, mock-data.csv and marketing-asset.svg.
5. If a link requests access or the version differs, stop and ask the trainer to provide the current course record; do not use archived copies.

Topic 1: Plan the campaign

01 Campaign objective and KPI tree

Inputs: business goal, audience, baseline. AI-assisted operation: map goal to one primary conversion event. Deliverable: SMART brief and KPI owner.

Decision rule: Choose lead volume only with quality guardrail. Common failure: Vanity engagement masks low-intent traffic.

Worked mock-data comparison (Leads): baseline 18; A 25; B 32. Check denominators and business relevance before selecting a variant.

Source: official; legacy: 11, 43–45. Companion slides: 6–10.

02 Audience evidence and consent

Inputs: aggregate insights, interview notes, consent status. AI-assisted operation: separate observed signal from AI inference. Deliverable: evidence-backed audience segment.

Decision rule: Reject personas without a source and confidence tag. Common failure: Inferred sensitive traits can breach trust.

Worked mock-data comparison (Qualified visits): baseline 210; A 260; B 305. Check denominators and business relevance before selecting a variant.

Source: IBM;legacy:38-42. Companion slides: 11–15.

03 Customer journey mapping

Inputs: awareness questions, comparison needs, conversion friction. AI-assisted operation: link each stage to content, CTA and metric. Deliverable: journey-to-content matrix.

Decision rule: Every stage needs a testable next action. Common failure: A single post cannot serve every intent.

Worked mock-data comparison (Click rate): baseline 1.5; A 2.1; B 2.8. Check denominators and business relevance before selecting a variant.

Source: legacy:46-51. Companion slides: 16–20.

04 Competitive content audit

Inputs: five public competitor posts, dates, formats. AI-assisted operation: code hooks, proof, CTA and engagement by format. Deliverable: gap and differentiation map.

Decision rule: Adapt the unmet need, not the competitor wording. Common failure: Copying creative weakens differentiation.

Worked mock-data comparison (Save rate): baseline 1.2; A 1.8; B 2.4. Check denominators and business relevance before selecting a variant.

Source: legacy:52-55. Companion slides: 21–25.

05 Campaign brief and brand voice

Inputs: offer facts, audience, exclusions, voice rules. AI-assisted operation: constrain AI output with proof and prohibited claims. Deliverable: approved creative brief.

Decision rule: Approve only claims supported by source facts. Common failure: AI may invent benefits and testimonials.

Worked mock-data comparison (Approval rate): baseline 42; A 65; B 81. Check denominators and business relevance before selecting a variant.

Source: HBR;legacy:56-58. Companion slides: 26–30.

06 Content pillar allocation

Inputs: brand goal, journey stages, available production hours. AI-assisted operation: allocate educate, proof and offer posts. Deliverable: four-week pillar mix.

Decision rule: Rebalance when a pillar misses its intended KPI. Common failure: Posting volume is not a strategy.

Worked mock-data comparison (Qualified actions): baseline 14; A 19; B 27. Check denominators and business relevance before selecting a variant.

Source: legacy:59-61. Companion slides: 31–35.

07 Prompt specification

Inputs: role, audience, facts, format, constraints. AI-assisted operation: supply structured context and acceptance criteria. Deliverable: repeatable prompt template.

Decision rule: Keep facts and output schema explicit. Common failure: Vague prompts produce generic copy.

Worked mock-data comparison (Usable drafts): baseline 2; A 5; B 8. Check denominators and business relevance before selecting a variant.

Source: AI Strategy 2025. Companion slides: 36–40.

08 Source-grounded claim review

Inputs: product sheet, price list, campaign brief. AI-assisted operation: trace every generated claim to approved evidence. Deliverable: claim-to-source register.

Decision rule: Hold unverified comparative or health claims. Common failure: Confident output is not proof.

Worked mock-data comparison (Claims approved): baseline 5; A 8; B 10. Check denominators and business relevance before selecting a variant.

Source: IBM;HBR. Companion slides: 41–45.

09 Community advocate plan

Inputs: customer feedback, permission log, loyalty goals. AI-assisted operation: identify advocates and design a voluntary contribution loop. Deliverable: advocate engagement plan.

Decision rule: Use consented stories with attribution. Common failure: Incentives can distort authentic reviews.

Worked mock-data comparison (Opt-in advocates): baseline 8; A 15; B 23. Check denominators and business relevance before selecting a variant.

Source: legacy:25,176-178. Companion slides: 46–50.

10 Campaign economics

Inputs: gross margin, conversion rate, media budget. AI-assisted operation: derive allowable CPA and break-even ROAS. Deliverable: budget guardrail sheet.

Decision rule: Stop spend above allowable CPA after sufficient evidence. Common failure: ROAS alone ignores margin and fulfillment.

Worked mock-data comparison (Contribution margin): baseline 20; A 28; B 36. Check denominators and business relevance before selecting a variant.

Source: legacy:163-174. Companion slides: 51–55.

Topic 2: Evaluate platforms and opportunity

11 Platform feasibility scoring

Inputs: audience fit, format fit, cost, capability. AI-assisted operation: weight factors and score channels consistently. Deliverable: platform selection matrix.

Decision rule: Pilot only channels above threshold and with owner. Common failure: Popularity is not audience fit.

Worked mock-data comparison (Fit score): baseline 58; A 72; B 83. Check denominators and business relevance before selecting a variant.

Source: legacy:18-24. Companion slides: 56–60.

12 Instagram format choice

Inputs: creative inventory, placement, audience behavior. AI-assisted operation: match Feed, Stories or Reels to message and asset. Deliverable: format-to-CTA map.

Decision rule: Use motion where demonstration adds evidence. Common failure: Old IGTV guidance is obsolete.

Worked mock-data comparison (Completion rate): baseline 12; A 19; B 27. Check denominators and business relevance before selecting a variant.

Source: legacy:67-95. Companion slides: 61–65.

13 Facebook and Instagram roles

Inputs: community intent, account assets, response capacity. AI-assisted operation: separate discovery and community service objectives. Deliverable: cross-placement plan.

Decision rule: Keep creative adapted to placement context. Common failure: Duplicated creative may crop or underperform.

Worked mock-data comparison (Response rate): baseline 35; A 49; B 62. Check denominators and business relevance before selecting a variant.

Source: legacy:98-104. Companion slides: 66–70.

14 LinkedIn B2B feasibility

Inputs: buying committee, professional proof, lead value. AI-assisted operation: score professional audience against acquisition cost. Deliverable: B2B pilot rationale.

Decision rule: Test only when deal economics support cost. Common failure: Broad consumer hooks may misfit B2B.

Worked mock-data comparison (Qualified leads): baseline 7; A 12; B 17. Check denominators and business relevance before selecting a variant.

Source: legacy:platforms. Companion slides: 71–75.

15 TikTok short-video feasibility

Inputs: demonstration potential, editing time, audience age. AI-assisted operation: compare hook retention and production burden. Deliverable: video pilot decision.

Decision rule: Require native-format proof before budget scale. Common failure: Trend imitation can weaken brand fit.

Worked mock-data comparison (3-second hold): baseline 24; A 35; B 46. Check denominators and business relevance before selecting a variant.

Source: official. Companion slides: 76–80.

16 Audience overlap and exclusions

Inputs: first-party segments, channel reach, frequency. AI-assisted operation: deduplicate target groups and set exclusions. Deliverable: channel reach plan.

Decision rule: Prioritize incremental reach over raw impressions. Common failure: Overlap inflates apparent scale.

Worked mock-data comparison (Unique reach): baseline 1200; A 1600; B 1950. Check denominators and business relevance before selecting a variant.

Source: legacy:137-139. Companion slides: 81–85.

17 Local market relevance

Inputs: Singapore language, price, cultural context. AI-assisted operation: review generated variants for local accuracy. Deliverable: localization QA sheet.

Decision rule: Reject unsupported local stereotypes. Common failure: Literal translation misses local context.

Worked mock-data comparison (Local relevance): baseline 61; A 76; B 88. Check denominators and business relevance before selecting a variant.

Source: Equinet;IJRSI. Companion slides: 86–90.

18 Opportunity sizing

Inputs: addressable segment, reachable share, response rate. AI-assisted operation: estimate attainable leads with explicit assumptions. Deliverable: funnel estimate.

Decision rule: Revise assumptions after pilot evidence. Common failure: AI estimates can appear more precise than data.

Worked mock-data comparison (Expected leads): baseline 20; A 34; B 46. Check denominators and business relevance before selecting a variant.

Source: legacy:152. Companion slides: 91–95.

19 Platform privacy fit

Inputs: data categories, pixel consent, vendor terms. AI-assisted operation: score data need against privacy and capability. Deliverable: privacy-aware platform decision.

Decision rule: Avoid activating unconsented personal data. Common failure: Ad targeting can expose private information.

Worked mock-data comparison (Consented events): baseline 40; A 65; B 86. Check denominators and business relevance before selecting a variant.

Source: legacy:127-131. Companion slides: 96–100.

20 Pilot test design

Inputs: hypothesis, audience, success threshold. AI-assisted operation: hold one variable while assigning equal budgets. Deliverable: channel pilot plan.

Decision rule: Scale only after threshold and sample check. Common failure: Multiple changes erase causal learning.

Worked mock-data comparison (Conversion rate): baseline 1.1; A 1.5; B 2. Check denominators and business relevance before selecting a variant.

Source: legacy:153-158. Companion slides: 101–105.

Topic 3: Create the campaign

21 AI-assisted caption writing

Inputs: approved offer, proof points, persona question. AI-assisted operation: generate hook, value, evidence and CTA variants. Deliverable: caption shortlist.

Decision rule: Human reviewer checks every fact and tone. Common failure: Invented urgency erodes trust.

Worked mock-data comparison (Click rate): baseline 1.2; A 1.9; B 2.5. Check denominators and business relevance before selecting a variant.

Source: legacy:85;Braze. Companion slides: 106–110.

22 Hashtag and location strategy

Inputs: topic taxonomy, local service area, platform policy. AI-assisted operation: choose descriptive tags with relevance tests. Deliverable: tag and location plan.

Decision rule: Keep only tags tied to intent and geography. Common failure: High-volume tags can attract wrong audience.

Worked mock-data comparison (Discovery visits): baseline 90; A 130; B 170. Check denominators and business relevance before selecting a variant.

Source: legacy:85-87. Companion slides: 111–115.

23 AI image brief

Inputs: product photo, usage rights, visual identity. AI-assisted operation: specify scene, composition and protected brand elements. Deliverable: image generation brief.

Decision rule: Review hands, product fidelity and disclosure need. Common failure: Generated products may misrepresent reality.

Worked mock-data comparison (Asset approval): baseline 45; A 65; B 82. Check denominators and business relevance before selecting a variant.

Source: IBM; legacy: 63-76. Companion slides: 116–120.

24 Short-video storyboard

Inputs: one customer problem, proof, CTA, duration. AI-assisted operation: sequence hook, demonstration, evidence and action. Deliverable: shot-by-shot storyboard.

Decision rule: Cut scenes that do not advance message. Common failure: Motion without proof produces weak retention.

Worked mock-data comparison (6-second hold): baseline 28; A 39; B 51. Check denominators and business relevance before selecting a variant.

Source: legacy: 90-96. Companion slides: 121–125.

25 Creative variant matrix

Inputs: one offer, three hooks, two visuals. AI-assisted operation: cross controlled variants and label each cell. Deliverable: six-cell test matrix.

Decision rule: Test one causal variable per comparison. Common failure: Confounded variants hide what worked.

Worked mock-data comparison (Qualified clicks): baseline 42; A 56; B 70. Check denominators and business relevance before selecting a variant.

Source: McKinsey; legacy: 114-115. Companion slides: 126–130.

26 Ad campaign hierarchy

Inputs: objective, audience, placement, creative. AI-assisted operation: map campaign to ad set to ad with ownership. Deliverable: editable hierarchy diagram.

Decision rule: Audit settings at each level before release. Common failure: Wrong objective optimizes the wrong event.

Worked mock-data comparison (Lead rate): baseline 1.4; A 2; B 2.6. Check denominators and business relevance before selecting a variant.

Source: legacy: 105-126. Companion slides: 131–135.

27 Budget and schedule controls

Inputs: daily cap, flight dates, business hours. AI-assisted operation: allocate spend by learning period and response capacity. Deliverable: flight and pacing plan.

Decision rule: Pause when spend exceeds guardrail without signals. Common failure: Automation can overspend poor creative.

Worked mock-data comparison (Cost per lead): baseline 34; A 28; B 22. Check denominators and business relevance before selecting a variant.

Source: legacy:163-175. Companion slides: 136–140.

28 Content calendar orchestration

Inputs: pillar mix, approvals, campaign dates. AI-assisted operation: schedule adapted assets with reviewer checkpoints. Deliverable: calendar and status board.

Decision rule: Publish only approved versions. Common failure: Auto-scheduling can propagate errors.

Worked mock-data comparison (On-time posts): baseline 55; A 72; B 90. Check denominators and business relevance before selecting a variant.

Source: legacy:159-160. Companion slides: 141–145.

29 UGC rights workflow

Inputs: customer post, consent record, attribution. AI-assisted operation: verify usage scope before repost or AI edit. Deliverable: rights and approval register.

Decision rule: Archive proof of permission and expiry. Common failure: Public visibility does not grant reuse rights.

Worked mock-data comparison (Approved UGC): baseline 4; A 7; B 10. Check denominators and business relevance before selecting a variant.

Source: legacy:127-129,176-178. Companion slides: 146–150.

30 Community response triage

Inputs: review sentiment, severity, order details. AI-assisted operation: classify reply, escalation and private resolution. Deliverable: response decision tree.

Decision rule: Human handles safety, legal and personal issues. Common failure: Bot replies can disclose account data.

Worked mock-data comparison (Resolved cases): baseline 14; A 21; B 28. Check denominators and business relevance before selecting a variant.

Source: legacy:176-178;Text.com. Companion slides: 151–155.

Topic 4: Measure and improve impact

31 Tracking event map

Inputs: landing page, click IDs, consent, lead form. AI-assisted operation: map ad click to qualified lead and CRM status. Deliverable: event and UTM schema.

Decision rule: Validate event firing with test conversion. Common failure: Pixel fires without consent or wrong event.

Worked mock-data comparison (Matched leads): baseline 35; A 52; B 68. Check denominators and business relevance before selecting a variant.

Source: legacy:133-135,150-151. Companion slides: 156–160.

32 Metric hierarchy

Inputs: reach, engagement, click, conversion, revenue. AI-assisted operation: separate diagnostic from outcome metrics. Deliverable: KPI dashboard model.

Decision rule: Judge campaign on objective metric. Common failure: Likes cannot prove commercial impact.

Worked mock-data comparison (Lead conversion): baseline 1.3; A 2.1; B 3. Check denominators and business relevance before selecting a variant.

Source: legacy:136-141. Companion slides: 161–165.

33 CTR and CPC diagnosis

Inputs: impressions, clicks, spend. AI-assisted operation: calculate CTR and CPC by ad variant. Deliverable: creative efficiency report.

Decision rule: Refresh hook when CTR falls with stable reach. Common failure: Low CPC can still bring poor leads.

Worked mock-data comparison (CTR): baseline 1.1; A 1.8; B 2.4. Check denominators and business relevance before selecting a variant.

Source: legacy:138-140. Companion slides: 166–170.

34 CPA and ROAS diagnosis

Inputs: spend, tracked conversions, revenue. AI-assisted operation: calculate CPA, ROAS and contribution return. Deliverable: unit economics table.

Decision rule: Compare with margin-adjusted break-even. Common failure: Attribution may omit offline sales.

Worked mock-data comparison (ROAS): baseline 1.4; A 2; B 2.8. Check denominators and business relevance before selecting a variant.

Source: legacy:141,168. Companion slides: 171–175.

35 Frequency and fatigue

Inputs: reach, impressions, weekly creative response. AI-assisted operation: calculate frequency and compare decay. Deliverable: fatigue alert.

Decision rule: Rotate creative when frequency rises and CTR falls. Common failure: More impressions may mean saturation.

Worked mock-data comparison (Frequency): baseline 1.8; A 2.6; B 3.7. Check denominators and business relevance before selecting a variant.

Source: legacy:137-138. Companion slides: 176–180.

36 A/B test interpretation

Inputs: two variants, same audience, sample size. AI-assisted operation: compare conversion and uncertainty before choosing. Deliverable: test decision memo.

Decision rule: Do not crown a winner on tiny samples. Common failure: Peeking can create false confidence.

Worked mock-data comparison (Conversion rate): baseline 1.2; A 1.6; B 2. Check denominators and business relevance before selecting a variant.

Source: legacy:153-162. Companion slides: 181–185.

37 AI-assisted report narrative

Inputs: clean export, metric definitions, anomalies. AI-assisted operation: ask AI to summarize changes with cited rows. Deliverable: reviewed weekly insight memo.

Decision rule: Every statement must trace to data cell. Common failure: AI can invent causality from correlation.

Worked mock-data comparison (Actionable insights): baseline 2; A 4; B 6. Check denominators and business relevance before selecting a variant.

Source: legacy:142-149;IBM. Companion slides: 186–190.

38 Retargeting eligibility

Inputs: consent, site visits, exclusions, window. AI-assisted operation: build eligible audience with suppression rules. Deliverable: retargeting policy.

Decision rule: Exclude converters and honor consent changes. Common failure: Persistent targeting can feel intrusive.

Worked mock-data comparison (Recovered leads): baseline 6; A 11; B 17. Check denominators and business relevance before selecting a variant.

Source: legacy:150-151. Companion slides: 191–195.

39 Budget reallocation

Inputs: segment CPA, lead quality, delivery constraints. AI-assisted operation: shift spend toward incremental qualified return. Deliverable: reallocation decision.

Decision rule: Reserve a test budget and cap downside. Common failure: Optimizing solely to ROAS may starve discovery.

Worked mock-data comparison (Qualified leads): baseline 14; A 20; B 27. Check denominators and business relevance before selecting a variant.

Source: legacy:167-174. Companion slides: 196–200.

40 Organisational impact review

Inputs: baseline effort, cost, quality, outcomes. AI-assisted operation: compare AI-assisted workflow with control and risk cost. Deliverable: impact scorecard.

Decision rule: Adopt only with measured quality and governance. Common failure: Speed gains can hide review burden.

Worked mock-data comparison (Hours saved): baseline 4; A 8; B 12. Check denominators and business relevance before selecting a variant.

Source: HBR;CMR;McKinsey. Companion slides: 201–205.

Detailed lab workflow

Lab 01: Campaign objective and KPI tree to Campaign brief and brand voice

1. Read the fictional product brief; list price, contribution margin, unsupported claims and campaign constraints.
2. In mock-data.csv, distinguish observed audience evidence from synthetic assumptions and mark confidence for each segment.
3. Write one business goal, one primary conversion event, baseline, target, owner and deadline.
4. Build a three-stage journey map with customer question, evidence and CTA for every stage.
5. Audit five synthetic competitor examples by hook, format, proof and CTA; record one unmet need.
6. Create a campaign brief with approved claims, prohibited claims, voice traits, target segment and review owner.
7. Save the KPI tree, journey map and reviewed brief.

Lab 02: Content pillar allocation to Campaign economics

1. Allocate a four-week content mix to educate, prove and offer; justify the share by journey need.
2. Use the prompt PDF to generate two drafts from the approved brief; compare each against the acceptance criteria.
3. Create a claim register with generated statement, source fact, reviewer and approval status.
4. Design a consent-based advocate request and a usage-rights record for customer material.
5. Compute allowable CPA from S\$14 contribution margin and explain why ROAS alone is insufficient.

6. Save the approved prompt, rejected draft, claim register and economics sheet.

Lab 03: Platform feasibility scoring to TikTok short-video feasibility

1. Set platform weights for audience fit, format fit, cost and team capacity; ensure weights sum to 100%.
2. Score Instagram, Facebook, LinkedIn and TikTok from 1–5 using written evidence and assumptions.
3. For Instagram, map Feed, Stories and Reels to the customer journey and asset requirement.
4. Specify one native-format creative adaptation for each shortlisted platform.
5. Estimate production effort and response capacity; compare with the available team hours.
6. Apply a threshold and select one pilot platform plus one reserve; save the scoring matrix.

Lab 04: Audience overlap and exclusions to Pilot test design

1. Estimate incremental unique reach after removing overlapping audience segments.
2. Review the Singapore-local draft for price, spelling, culture and unsupported stereotypes.
3. Calculate an opportunity funnel from reachable audience, response rate and lead qualification assumptions.
4. Identify personal-data categories in the draft tracking plan; mark the consent and vendor-control needed for each.
5. State a one-variable pilot hypothesis, test period, equal-budget rule and decision threshold.
6. Save the funnel, privacy gate and pilot approval memo.

Lab 05: AI-assisted caption writing to Creative variant matrix

1. Write three caption hooks from the approved product sheet; add benefit, evidence and CTA to each.
2. Select hashtags by topic relevance and geography, recording why each tag fits the audience.
3. Use the editable SVG marketing asset as a wireframe; define a separate AI image brief with product-fidelity constraints.
4. Storyboard a 15-second video: hook, demonstration, evidence and CTA with time per shot.
5. Build a six-cell test matrix from three hooks and two visuals; label the one variable compared in each test.

6. Reject any generated claim or visual that misrepresents the fictional product; save corrected assets.

Lab 06: Ad campaign hierarchy to Community response triage

1. Diagram campaign, ad set and ad levels and place objective, audience, placement, budget and creative at the right level.
2. Set a synthetic daily budget, flight dates, frequency check and stop-spend guardrail.
3. Populate a four-week editorial calendar with owner, asset status and review date.
4. For a mock customer post, record consent scope, attribution and expiry before reuse.
5. Classify three mock comments as reply, private resolution or escalation; draft a human-reviewed response.
6. Save hierarchy, calendar, rights log and response decision tree without publishing a live ad.

Lab 07: Tracking event map to Frequency and fatigue

1. Create a UTM naming convention and map ad click, landing view, lead submit and qualified lead.
2. Using mock-data.csv, calculate $CTR = \text{clicks/impressions}$ and $CPC = \text{spend/clicks}$; state units.
3. Calculate $CPA = \text{spend/qualified leads}$ and $ROAS = \text{attributed revenue/spend}$; compare with margin.
4. Calculate frequency = impressions/reach; compare rising frequency with declining CTR to diagnose fatigue.
5. Check consent before any pixel/event plan; validate a synthetic test conversion and avoid real customer data.
6. Save the event map, formula sheet and a data-quality note.

Lab 08: A/B test interpretation to Organisational impact review

1. Write one A/B hypothesis, hold one variable constant and set a sample-size caveat.
2. Compare the two variants on primary conversion metric and lead quality, avoiding a decision from tiny samples.
3. Prompt AI to draft a weekly narrative from the synthetic rows; trace every claim to a cell.
4. Specify retargeting eligibility, consent window and converter suppression.
5. Allocate a small next-test budget using qualified return and a downside cap.
6. Compare baseline versus AI-assisted hours, approval quality, media outcome and review burden; save an impact memo.

Measurement and governance

CTR = clicks ÷ impressions × 100. CPC = spend ÷ clicks. CPA = spend ÷ qualified conversions. ROAS = attributed revenue ÷ spend. Conversion rate = qualified conversions ÷ clicks × 100. Compare against margin and attribution limits, not just top-line engagement.

Obtain consent for customer data and UGC; respect image and text rights; review AI outputs for factual accuracy, local relevance and bias; keep a human approval record. Use platform tracking only when the organisation has approved its privacy and consent process.

References

official: <https://www.tertiarycourses.com.sg/generative-ai-for-social-media-marketing.html>

IBM: <https://www.ibm.com/think/topics/generative-ai-marketing>

McKinsey: <https://www.mckinsey.com/capabilities/growth-marketing-and-sales/our-insights/how-generative-ai-can-boost-consumer-marketing>

HBR: <https://hbr.org/2025/03/how-should-gen-ai-fit-into-your-marketing-strategy>

CMR: <https://cmr.berkeley.edu/2025/11/from-novelty-to-autopilot-how-generative-ai-is-reshaping-marketing/>

Braze: <https://www.braze.com/resources/articles/generative-ai-for-marketing>

Text.com: <https://www.text.com/blog/generative-ai-in-marketing/>

Equinet: <https://www.equinetacademy.com/blog/generative-ai-for-digital-marketing-a-practical-guide/>

IJRSI: <https://rsisinternational.org/journals/ijrsi/articles/the-impact-of-generative-ai-on-digital-marketing-strategies-evidence-from-emerging-markets/>

TechTarget: <https://www.techtarget.com/enterprise-software/feature/How-to-use-generative-AI-for-marketing>

Intuz: <https://www.intuz.com/blog/generative-ai-in-digital-marketing/>